




BITS Pilani
Pilani | Dubai | Goa | Hyderabad

DEEP NEURAL NETWORK

MODULE # 5: DFNN NUMERICAL

DNN Team
BITS Pilani WILP

TABLE OF CONTENTS

- 1 COMPLETE NEURAL NETWORK TRAINING
- 2 GRADIENT FLOW
- 3 DESIGN A NEURAL NETWORK
- 4 ACTIVATION FUNCTIONS
- 5 LOSS FUNCTIONS AND THEIR DERIVATIVES
- 6 COMPARISON AND SUMMARY



NUMERICAL PROBLEM : FORWARD PROPAGATION

Problem: Consider a simple 2-layer neural network with the following parameters. Calculate the forward propagation step by step.


Network Architecture:

- Input layer: 2 neurons
- Hidden layer: 3 neurons (ReLU activation)
- Output layer: 2 neurons (Sigmoid activation)

Given Parameters:

$$\checkmark x = \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix} \quad W^{[1]} = \begin{bmatrix} 0.5 & 0.2 \\ -0.3 & 0.8 \\ 0.1 & -0.4 \end{bmatrix}, \quad b^{[1]} = \begin{bmatrix} 0.1 \\ -0.2 \\ 0.3 \end{bmatrix}$$
$$W^{[2]} = \begin{bmatrix} 0.4 & -0.1 & 0.6 \\ 0.2 & 0.7 & -0.3 \end{bmatrix}, \quad b^{[2]} = \begin{bmatrix} 0.2 \\ -0.1 \end{bmatrix}$$

FORWARD PROP: STEP 0 - INITIAL NETWORK



Ready to compute forward propagation!

SOLUTION: FORWARD PROP - LAYER 1

Step 1: Calculate $z^{[1]}$

$$z^{[1]} = \underline{W^{[1]}}x + \underline{b^{[1]}}$$


$$z^{[1]} = \begin{bmatrix} 0.5 & 0.2 \\ -0.3 & 0.8 \\ 0.1 & -0.4 \end{bmatrix} \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix} + \begin{bmatrix} 0.1 \\ -0.2 \\ 0.3 \end{bmatrix}$$

$$= \begin{bmatrix} 0.5(1.0) + 0.2(0.5) \\ -0.3(1.0) + 0.8(0.5) \\ 0.1(1.0) + (-0.4)(0.5) \end{bmatrix} + \begin{bmatrix} 0.1 \\ -0.2 \\ 0.3 \end{bmatrix}$$

$$= \begin{bmatrix} 0.6 \\ 0.1 \\ -0.1 \end{bmatrix} + \begin{bmatrix} 0.1 \\ -0.2 \\ 0.3 \end{bmatrix} = \begin{bmatrix} 0.7 \\ -0.1 \\ 0.2 \end{bmatrix} \checkmark \rightarrow$$

$$\begin{bmatrix} 0.7 \\ 0.1 \\ 0.2 \end{bmatrix}$$

FORWARD PROP: STEP 1 - COMPUTING $z^{[1]}$



$$z^{[1]} = W^{[1]}x + b^{[1]} = \begin{bmatrix} 0.5 & 0.2 \\ -0.3 & 0.8 \\ 0.1 & -0.4 \end{bmatrix} \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix} + \begin{bmatrix} 0.1 \\ -0.2 \\ 0.3 \end{bmatrix} = \begin{bmatrix} 0.7 \\ -0.1 \\ 0.2 \end{bmatrix}$$

SOLUTION: FORWARD PROP - LAYER 1 ACTIVATION

Step 2: Apply ReLU activation $a^{[1]} = \text{ReLU}(z^{[1]})$

ReLU function: $\text{ReLU}(x) = \max(0, x)$


$$a^{[1]} = \text{ReLU} \begin{bmatrix} 0.7 \\ -0.1 \\ 0.2 \end{bmatrix} = \begin{bmatrix} \max(0, 0.7) \\ \max(0, -0.1) \\ \max(0, 0.2) \end{bmatrix} = \begin{bmatrix} 0.7 \\ 0.0 \\ 0.2 \end{bmatrix}$$

Hidden layer activation: $a^{[1]} = \begin{bmatrix} 0.7 \\ 0.0 \\ 0.2 \end{bmatrix}$

FORWARD PROP: STEP 2 - APPLYING RELU

Hidden Layer: ReLU Activation

Input Layer



$$a^{[1]} = \text{ReLU} \begin{bmatrix} 0.7 \\ -0.1 \\ 0.2 \end{bmatrix} = \begin{bmatrix} 0.7 \\ 0.0 \\ 0.2 \end{bmatrix}$$


Neuron 2 is **dead** ($z < 0$)

SOLUTION: FORWARD PROP - LAYER 2

Step 3: Calculate $z^{[2]}$

$$\begin{aligned} z^{[2]} &= \underline{W}^{[2]} \underline{a}^{[1]} + \underline{b}^{[2]} \\ z^{[2]} &= \begin{bmatrix} 0.4 & -0.1 & 0.6 \\ 0.2 & 0.7 & -0.3 \end{bmatrix} \begin{bmatrix} 0.7 \\ 0.0 \\ 0.2 \end{bmatrix} + \begin{bmatrix} 0.2 \\ -0.1 \end{bmatrix} \\ &= \begin{bmatrix} 0.4(0.7) + (-0.1)(0.0) + 0.6(0.2) \\ 0.2(0.7) + 0.7(0.0) + (-0.3)(0.2) \end{bmatrix} + \begin{bmatrix} 0.2 \\ -0.1 \end{bmatrix} \\ &= \begin{bmatrix} 0.28 + 0 + 0.12 \\ 0.14 + 0 - 0.06 \end{bmatrix} + \begin{bmatrix} 0.2 \\ -0.1 \end{bmatrix} = \begin{bmatrix} \underline{0.4} + \underline{0.2} \\ \underline{0.08} - \underline{0.1} \end{bmatrix} = \begin{bmatrix} 0.6 \\ -0.02 \end{bmatrix} \end{aligned}$$

FORWARD PROP: STEP 3 - COMPUTING $z^{[2]}$



$$z^{[2]} = W^{[2]}a^{[1]} + b^{[2]} = \begin{bmatrix} 0.4 & -0.1 & 0.6 \\ 0.2 & 0.7 & -0.3 \end{bmatrix} \begin{bmatrix} 0.7 \\ 0.0 \\ 0.2 \end{bmatrix} + \begin{bmatrix} 0.2 \\ -0.1 \end{bmatrix} = \begin{bmatrix} 0.6 \\ -0.02 \end{bmatrix}$$

SOLUTION: FORWARD PROP - OUTPUT LAYER

Step 4: Apply Sigmoid activation $a^{[2]} = \sigma(z^{[2]})$

Sigmoid function: $\sigma(x) = \frac{1}{1 + e^{-x}}$


$$a^{[2]} = \sigma \begin{bmatrix} 0.6 \\ -0.02 \end{bmatrix}$$

$$= \begin{bmatrix} \frac{1}{1+e^{-0.6}} \\ \frac{1}{1+e^{0.02}} \end{bmatrix}$$

Output activation: $a^{[2]} = \begin{bmatrix} 0.646 \\ 0.495 \end{bmatrix}$




FORWARD PROP: STEP 4 - APPLYING SIGMOID



Forward propagation complete! Prediction: $\hat{y} = [0.646, 0.495]^T$

COMPUTATIONAL GRAPH: FORWARD PROPAGATION



NUMERICAL PROBLEM: BACKWARD PROPAGATION

Problem: Using the previous network, calculate the backward propagation.

Given:

- Forward propagation results from previous problem
- Target output: $y = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$
- Loss function: Binary Cross-Entropy (BCE)

Calculate the gradients $\frac{\partial L}{\partial W^{[2]}}$, $\frac{\partial L}{\partial b^{[2]}}$, $\frac{\partial L}{\partial W^{[1]}}$, and $\frac{\partial L}{\partial b^{[1]}}$.

SOLUTION: BACKWARD PROPAGATION - LOSS


Step 1: Calculate the Binary Cross-Entropy loss

For our outputs: $a^{[2]} = \begin{bmatrix} 0.646 \\ 0.495 \end{bmatrix}$ $y = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$

BCE Loss Formula:
$$L = -\frac{1}{n} \sum_{i=1}^n [y_i \log(a_i^{[2]}) + (1 - y_i) \log(1 - a_i^{[2]})]$$

$$\begin{aligned} L &= -\frac{1}{2} [1 \cdot \log(0.646) + (1 - 1) \cdot \log(1 - 0.646) \\ &\quad + 0 \cdot \log(0.495) + (1 - 0) \cdot \log(1 - 0.495)] \\ &= -\frac{1}{2} [\log(0.646) + \log(0.505)] \\ &= -\frac{1}{2} [-0.437 + (-0.683)] = 0.560 \end{aligned}$$

COMPUTING LOSS



$$L = -\frac{1}{2} [y_1 \log(\hat{y}_1) + (1 - y_1) \log(1 - \hat{y}_1) + y_2 \log(\hat{y}_2) + (1 - y_2) \log(1 - \hat{y}_2)]$$
$$L = -\frac{1}{2} [\log(0.646) + \log(0.505)] = -\frac{1}{2} [-0.437 - 0.683] = 0.560$$


Now we backpropagate to update weights!

SOLUTION: BACKPROP - OUTPUT LAYER ERROR

Step 2: Calculate output layer error $\delta^{[2]}$

$$\frac{\partial L}{\partial z^{[2]}} = \delta^{[2]} = a^{[2]} - y$$
$$\delta^{[2]} = \begin{bmatrix} 0.646 \\ 0.495 \end{bmatrix} - \begin{bmatrix} 1 \\ 0 \end{bmatrix} = \begin{bmatrix} -0.354 \\ 0.495 \end{bmatrix}$$

BACK-PROP STEP 1 - OUTPUT LAYER ERROR



$$\delta^{[2]} = \begin{bmatrix} 0.646 \\ 0.495 \end{bmatrix} - \begin{bmatrix} 1 \\ 0 \end{bmatrix} = \begin{bmatrix} -0.354 \\ 0.495 \end{bmatrix}$$


SOLUTION: BACKWARD PROP - OUTPUT LAYER

Step 3: Calculate gradients for output layer

Recall: $a^{[1]} = [0.7 \ 0.0 \ 0.2]^T$

$$\begin{aligned}\frac{\partial L}{\partial W^{[2]}} &= \delta^{[2]} (a^{[1]})^T = \begin{bmatrix} -0.354 \\ 0.495 \end{bmatrix} \begin{bmatrix} 0.7 & 0.0 & 0.2 \end{bmatrix} \\ &= \begin{bmatrix} -0.354 \times 0.7 & -0.354 \times 0.0 & -0.354 \times 0.2 \\ 0.495 \times 0.7 & 0.495 \times 0.0 & 0.495 \times 0.2 \end{bmatrix} \\ &= \begin{bmatrix} -0.248 & 0.000 & -0.071 \\ 0.347 & 0.000 & 0.099 \end{bmatrix} \\ \frac{\partial L}{\partial b^{[2]}} &= \delta^{[2]} = \begin{bmatrix} -0.354 \\ 0.495 \end{bmatrix}\end{aligned}$$

BACK-PROP: STEP 2 - OUTPUT LAYER GRADIENT




$$\frac{\partial L}{\partial w^{[2]}} = \delta^{[2]} (a^{[1]})^T = \begin{bmatrix} -0.354 \\ 0.495 \end{bmatrix} \begin{bmatrix} 0.7 & 0.0 & 0.2 \end{bmatrix} = \begin{bmatrix} -0.248 & 0.000 & -0.071 \\ 0.347 & 0.000 & 0.099 \end{bmatrix}$$

$$\frac{\partial L}{\partial b^{[2]}} = \delta^{[2]} = \begin{bmatrix} -0.354 \\ 0.495 \end{bmatrix}$$

SOLUTION: BACKPROP - HIDDEN LAYER ERROR


Step 4: Calculate hidden layer error $\delta^{[1]}$

Backpropagate error to hidden layer: $\delta^{[1]} = (W^{[2]})^T \delta^{[2]} \odot g'^{[1]}(z^{[1]})$


$$\begin{aligned}(W^{[2]})^T \delta^{[2]} &= \begin{bmatrix} 0.4 & 0.2 \\ -0.1 & 0.7 \\ 0.6 & -0.3 \end{bmatrix} \begin{bmatrix} -0.354 \\ 0.495 \end{bmatrix} \\ &= \begin{bmatrix} 0.4(-0.354) + 0.2(0.495) \\ -0.1(-0.354) + 0.7(0.495) \\ 0.6(-0.354) + (-0.3)(0.495) \end{bmatrix} \\ &= \begin{bmatrix} -0.142 + 0.099 \\ 0.035 + 0.347 \\ -0.212 - 0.149 \end{bmatrix} = \begin{bmatrix} -0.043 \\ 0.382 \\ -0.361 \end{bmatrix}\end{aligned}$$

BACK-PROP: STEP 3 - BACKPROP TO HIDDEN LAYERS

Backprop: $(W^{[2]})^T \delta^{[2]}$



$$(W^{[2]})^T \delta^{[2]} = \begin{bmatrix} 0.4 & 0.2 \\ -0.1 & 0.7 \\ 0.6 & -0.3 \end{bmatrix} \begin{bmatrix} -0.354 \\ 0.495 \end{bmatrix} = \begin{bmatrix} -0.043 \\ 0.382 \\ -0.361 \end{bmatrix}$$

SOLUTION: BACKPROP - HIDDEN LAYER ERROR

Step 4 (continued):


$$g'^{[1]}(z^{[1]}) = \begin{bmatrix} 1 & (\text{since } z_1^{[1]} = 0.7 > 0) \\ 0 & (\text{since } z_2^{[1]} = -0.1 \leq 0) \\ 1 & (\text{since } z_3^{[1]} = 0.2 > 0) \end{bmatrix}$$

Error to Layer 1: $\delta^{[1]} = (W^{[2]})^T \delta^{[2]} \odot g'^{[1]}(z^{[1]})$

$$\delta^{[1]} = \begin{bmatrix} -0.043 \\ 0.382 \\ -0.361 \end{bmatrix} \odot \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix} = \begin{bmatrix} -0.043 \\ 0.000 \\ -0.361 \end{bmatrix}$$

BACK-PROP: STEP 4 - APPLY RELU DERIVATIVE

Apply ReLU': $\delta^{[1]} = (W^{[2]})^T \delta^{[2]} \odot g'(z^{[1]})$



$$\delta^{[1]} = \begin{bmatrix} -0.043 \\ 0.382 \\ -0.361 \end{bmatrix} \odot \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix} = \begin{bmatrix} -0.043 \\ 0.0 \\ -0.361 \end{bmatrix}$$

Dead neuron gradient = 0 (no update!)

SOLUTION: BACKPROP - HIDDEN LAYER GRADIENT


Step 5: Calculate gradients for hidden layer

$$\begin{aligned}\frac{\partial L}{\partial \mathbf{W}^{[1]}} &= \delta^{[1]} \mathbf{x}^T = \begin{bmatrix} -0.043 \\ 0.000 \\ -0.361 \end{bmatrix} \begin{bmatrix} 1.0 & 0.5 \end{bmatrix} \\ \frac{\partial L}{\partial \mathbf{W}^{[1]}} &= \begin{bmatrix} -0.043 \times 1.0 & -0.043 \times 0.5 \\ 0.000 \times 1.0 & 0.000 \times 0.5 \\ -0.361 \times 1.0 & -0.361 \times 0.5 \end{bmatrix} \\ &= \begin{bmatrix} -0.043 & -0.022 \\ 0.000 & 0.000 \\ -0.361 & -0.181 \end{bmatrix} \\ \frac{\partial L}{\partial \mathbf{b}^{[1]}} &= \delta^{[1]} = \begin{bmatrix} -0.043 \\ 0.000 \\ -0.361 \end{bmatrix}\end{aligned}$$

BACK-PROP: STEP 5 - HIDDEN LAYER GRADIENT

Computing $\frac{\partial L}{\partial W^{[1]}}$


Input



$$\frac{\partial L}{\partial W^{[1]}} = \delta^{[1]} \mathbf{x}^T = \begin{bmatrix} -0.043 \\ 0.0 \\ -0.361 \end{bmatrix} \begin{bmatrix} 1.0 & 0.5 \end{bmatrix} = \begin{bmatrix} -0.043 & -0.022 \\ 0.0 & 0.0 \\ -0.361 & -0.181 \end{bmatrix}$$

$$\frac{\partial L}{\partial b^{[1]}} = \delta^{[1]} = \begin{bmatrix} -0.043 \\ 0.0 \\ -0.361 \end{bmatrix}$$

COMPUTATIONAL GRAPH: BACKWARD PROPAGAT



NUMERICAL PROBLEM: WEIGHT UPDATE

Problem: Using the gradients computed in previous problem, update the network parameters with learning rate $\alpha = 0.1$.

Given gradients:

$$\begin{aligned}\frac{\partial L}{\partial W^{[2]}} &= \begin{bmatrix} -0.248 & 0.000 & -0.071 \\ 0.347 & 0.000 & 0.099 \end{bmatrix} & \frac{\partial L}{\partial b^{[2]}} &= \begin{bmatrix} -0.354 \\ 0.495 \end{bmatrix} \\ \frac{\partial L}{\partial W^{[1]}} &= \begin{bmatrix} -0.043 & -0.022 \\ 0.000 & 0.000 \\ -0.361 & -0.181 \end{bmatrix} & \frac{\partial L}{\partial b^{[1]}} &= \begin{bmatrix} -0.043 \\ 0.000 \\ -0.361 \end{bmatrix}\end{aligned}$$

Calculate the updated parameters.

SOLUTION: PARAMETER UPDATES

Update rule: $\theta_{\text{new}} = \theta_{\text{old}} - \alpha \frac{\partial L}{\partial \theta}$

$$\begin{aligned} W_{\text{new}}^{[2]} &= W^{[2]} - 0.1 \times \frac{\partial L}{\partial W^{[2]}} \\ W_{\text{new}}^{[2]} &= \begin{bmatrix} 0.4 & -0.1 & 0.6 \\ 0.2 & 0.7 & -0.3 \end{bmatrix} - 0.1 \times \begin{bmatrix} -0.248 & 0.000 & -0.071 \\ 0.347 & 0.000 & 0.099 \end{bmatrix} \\ &= \begin{bmatrix} 0.4 + 0.025 & -0.1 + 0.000 & 0.6 + 0.007 \\ 0.2 - 0.035 & 0.7 + 0.000 & -0.3 - 0.010 \end{bmatrix} \\ &= \begin{bmatrix} 0.425 & -0.100 & 0.607 \\ 0.165 & 0.700 & -0.310 \end{bmatrix} \\ b_{\text{new}}^{[2]} &= \begin{bmatrix} 0.2 \\ -0.1 \end{bmatrix} - 0.1 \times \begin{bmatrix} -0.354 \\ 0.495 \end{bmatrix} = \begin{bmatrix} 0.235 \\ -0.150 \end{bmatrix} \end{aligned}$$

SOLUTION: PARAMETER UPDATES (CONTINUED)

$$W_{\text{new}}^{[1]} = W^{[1]} - 0.1 \times \frac{\partial L}{\partial W^{[1]}}$$


$$W_{\text{new}}^{[1]} = \begin{bmatrix} 0.5 & 0.2 \\ -0.3 & 0.8 \\ 0.1 & -0.4 \end{bmatrix} - 0.1 \times \begin{bmatrix} -0.043 & -0.022 \\ 0.000 & 0.000 \\ -0.361 & -0.181 \end{bmatrix}$$

$$= \begin{bmatrix} 0.5 + 0.004 & 0.2 + 0.002 \\ -0.3 + 0.000 & 0.8 + 0.000 \\ 0.1 + 0.036 & -0.4 + 0.018 \end{bmatrix} = \begin{bmatrix} 0.504 & 0.202 \\ -0.300 & 0.800 \\ 0.136 & -0.382 \end{bmatrix}$$

$$b_{\text{new}}^{[1]} = \begin{bmatrix} 0.1 \\ -0.2 \\ 0.3 \end{bmatrix} - 0.1 \times \begin{bmatrix} -0.043 \\ 0.000 \\ -0.361 \end{bmatrix} = \begin{bmatrix} 0.104 \\ -0.200 \\ 0.336 \end{bmatrix}$$

COMPUTATIONAL GRAPH: WEIGHT UPDATE STEP

Layer 2 Updates



Layer 1 Updates




TABLE OF CONTENTS

- 1 COMPLETE NEURAL NETWORK TRAINING
- 2 GRADIENT FLOW
- 3 DESIGN A NEURAL NETWORK
- 4 ACTIVATION FUNCTIONS
- 5 LOSS FUNCTIONS AND THEIR DERIVATIVES
- 6 COMPARISON AND SUMMARY



NUMERICAL PROBLEM: GRADIENT FLOW ANALYSIS

Problem: Consider a 4-layer network with sigmoid activations. Analyze gradient flow through the network.

Given:

- All weight matrices have spectral norm $\|W^{[l]}\| = 0.8$
- Sigmoid derivatives at each layer: $g'^{[l]} = 0.25$ (maximum sigmoid derivative)
- Initial gradient magnitude at output: $|\delta^{[4]}| = 1.0$

Calculate the gradient magnitude at each layer and discuss the implications.

Gradient Propagation Formula:

$$|\delta^{[l-1]}| = |\delta^{[l]}| \cdot \|W^{[l]}\| \cdot |g'^{[l-1]}|$$

SOLUTION: GRADIENT FLOW ANALYSIS

Layer-by-layer gradient magnitude:

Layer 4 (Output): $|\delta^{[4]}| = 1.0$ (given)

Layer 3: $|\delta^{[3]}| = |\delta^{[4]}| \cdot \|W^{[4]}\| \cdot |g'^{[3]}| = 1.0 \times 0.8 \times 0.25 = 0.2$

Layer 2: $|\delta^{[2]}| = |\delta^{[3]}| \cdot \|W^{[3]}\| \cdot |g'^{[2]}| = 0.2 \times 0.8 \times 0.25 = 0.04$

Layer 1: $|\delta^{[1]}| = |\delta^{[2]}| \cdot \|W^{[2]}\| \cdot |g'^{[1]}| = 0.04 \times 0.8 \times 0.25 = 0.008$

Analysis:

- Gradient magnitude decreases by factor of $0.8 \times 0.25 = 0.2$ per layer
- After 3 backward steps: gradient is reduced to 0.8% of original
- First layer receives extremely weak learning signal
- This demonstrates the vanishing gradient problem in deep sigmoid networks

TABLE OF CONTENTS

- 1 COMPLETE NEURAL NETWORK TRAINING
- 2 GRADIENT FLOW
- 3 DESIGN A NEURAL NETWORK**
- 4 ACTIVATION FUNCTIONS
- 5 LOSS FUNCTIONS AND THEIR DERIVATIVES
- 6 COMPARISON AND SUMMARY



EXAMPLE 1: SMALL TABULAR DATASET

Problem Setup: Compute the number of layers and number of neurons in each layer

- $d = 20$ features (customer data)
- $M = 800$ training samples
- $c = 3$ classes (customer segments)

Equations to be used are

$$L \leq \log_2(M/(d \times c)) + 1$$

$$\text{Parameters} \leq M/5$$

EXAMPLE 1: SMALL TABULAR DATASET

- Step 1: Apply constraints

$$L \leq \log_2(800/(20 \times 3)) + 1 = \log_2(13.33) + 1 \approx 4.7 \Rightarrow L \leq 4$$

- Step 2: Choose starting architecture Since $M < 1000$, start with $L = 2$:
[20, 40, 3]

- Step 3: Parameter count check Parameters

$$= 20 \times 40 + 40 + 40 \times 3 + 3 = 800 + 40 + 120 + 3 = 963$$

Constraint: $963 \leq 800/5 = 160$ ✗ Too many parameters!

- Step 4: Reduce complexity

- ▶ Try [20, 20, 3]: Parameters = $20 \times 20 + 20 + 20 \times 3 + 3 = 483$ ✗ Still too many!
- ▶ Try [20, 10, 3]: Parameters = $20 \times 10 + 10 + 10 \times 3 + 3 = 243$ ✗
- ▶ Try [20, 6, 3]: Parameters = $20 \times 6 + 6 + 6 \times 3 + 3 = 147$ ✓

- Final recommendation: [20, 6, 3] with regularization

EXAMPLE 2: MEDIUM IMAGE DATASET

Problem Setup: Compute the number of layers and number of neurons in each layer

- $d = 784$ features (28×28 grayscale images)
- $M = 5000$ training samples
- $c = 10$ classes (digit classification)

Equations to be used are

$$L \leq \log_2(M/(d \times c)) + 1$$

$$\text{Parameters} \leq M/5$$

EXAMPLE 2: MEDIUM IMAGE DATASET

- Step 1: Relaxed approach

For medium datasets, start with $L = 3$: $[784, 2352, 784, 10]$

- Step 2: Parameter count check

Parameters = $784 \times 2352 + 2352 + 2352 \times 784 + 784 + 784 \times 10 + 10 \approx 3.7M$

Constraint $5000/5 = 1000$ ✗ Way too many!

- Step 3: Practical architecture

- ▶ Try $[784, 128, 64, 10]$:

Parameters = $784 \times 128 + 128 + 128 \times 64 + 64 + 64 \times 10 + 10 = 109,000$ ✗

- ▶ Try $[784, 64, 32, 10]$:

Parameters = $784 \times 64 + 64 + 64 \times 32 + 32 + 32 \times 10 + 10 = 52,458$ ✗

- ▶ Try $[784, 32, 10]$: Parameters = $784 \times 32 + 32 + 32 \times 10 + 10 = 25,450$ ✗

- ▶ Try $[784, 16, 10]$: Parameters = $784 \times 16 + 16 + 16 \times 10 + 10 = 12,714$ ✗

- Reality check: Need to use regularization techniques (dropout, L2) with larger architectures.

EXAMPLE 3: LARGE TEXT DATASET

Problem Setup: Compute the number of layers and number of neurons in each layer

- $d = 1000$ features (TF-IDF vectors)
- $M = 50000$ training samples
- $c = 5$ classes (sentiment categories)

Equations to be used are

$$L \leq \log_2(M/(d \times c)) + 1$$

$$\text{Parameters} \leq M/5$$

EXAMPLE 3: LARGE TEXT DATASET

- Step 1: Apply constraints

$$L \leq \log_2(50000/(1000 \times 5)) + 1 = \log_2(10) + 1 \approx 4.3 \Rightarrow L \leq 4$$

- Step 2: Choose architecture

Since $M \geq 10000$, can use $L = 4$: $[1000, 5000, 2000, 1000, 5]$

- Step 3: Parameter count check Parameters =

$$1000 \times 5000 + 5000 + 5000 \times 2000 + 2000 + 2000 \times 1000 + 1000 + 1000 \times 5 + 5 = 5M + 10M + 2M + 5K = 17M$$

Constraint: $17M \leq 50000/5 = 10000$ ✗ Way too many!

- Step 4: Reasonable architecture

Try $[1000, 512, 256, 128, 5]$:

$$\begin{aligned} \text{Parameters} &= 1000 \times 512 + 512 + 512 \times 256 + 256 + 256 \times 128 + 128 + 128 \times 5 + 5 \\ &= 512K + 131K + 32K + 640 = 675K \text{ parameters } \checkmark \end{aligned}$$

This is $675000/50000 = 13.5$ parameters per sample, which is reasonable.

- Final recommendation: $[1000, 512, 256, 128, 5]$ with batch normalization and dropout.

Additional Reading




TABLE OF CONTENTS

- 1 COMPLETE NEURAL NETWORK TRAINING
- 2 GRADIENT FLOW
- 3 DESIGN A NEURAL NETWORK
- 4 ACTIVATION FUNCTIONS**
- 5 LOSS FUNCTIONS AND THEIR DERIVATIVES
- 6 COMPARISON AND SUMMARY



SIGMOID FUNCTION: DEFINITION AND PROPERTIES

The sigmoid function maps any real number to the interval $(0, 1)$.

Definition:

$$\sigma(z) = \frac{1}{1 + e^{-z}} \quad (1)$$

Properties:

- Output range: $(0, 1)$
- Smooth and differentiable everywhere
- $\sigma(0) = 0.5$
- $\sigma(-z) = 1 - \sigma(z)$ (symmetry)
- As $z \rightarrow \infty$: $\sigma(z) \rightarrow 1$
- As $z \rightarrow -\infty$: $\sigma(z) \rightarrow 0$

Use cases:

- Binary classification (output layer)
- Gate mechanisms in LSTMs

Key values:

- $\sigma(-2) \approx 0.12$
- $\sigma(0) = 0.5$
- $\sigma(2) \approx 0.88$

SIGMOID FUNCTION: DERIVATIVE DERIVATION

Goal: Find $\frac{d\sigma}{dz}$ where $\sigma(z) = \frac{1}{1+e^{-z}}$

Step 1: Rewrite and apply chain rule

$$\sigma(z) = (1 + e^{-z})^{-1} \Rightarrow \frac{d\sigma}{dz} = -(1 + e^{-z})^{-2} \cdot (-e^{-z}) = \frac{e^{-z}}{(1 + e^{-z})^2} \quad (2)$$

Step 2: Factor to express in terms of $\sigma(z)$



$$\frac{d\sigma}{dz} = \frac{e^{-z}}{(1 + e^{-z})^2} = \frac{1}{1 + e^{-z}} \cdot \frac{e^{-z}}{1 + e^{-z}} \quad (3)$$

Step 3: Recognize that $\sigma(z) = \frac{1}{1+e^{-z}}$ and $1 - \sigma(z) = \frac{e^{-z}}{1+e^{-z}}$

$$\frac{d\sigma}{dz} = \sigma(z) \cdot (1 - \sigma(z)) \quad (4)$$

Beautiful property: Derivative computed using only the function value itself!

SIGMOID FUNCTION: VISUALIZATION



Key observations:

- Sigmoid squashes inputs to $(0, 1)$ - ideal for probabilities
- Derivative peaks at $z = 0$ with value 0.25
- Derivative vanishes for large $|z|$ (vanishing gradient problem!)

SOFTMAX FUNCTION: DEFINITION

The softmax function converts a vector of real numbers into a probability distribution.

For input vector $\mathbf{z} = [z_1, z_2, \dots, z_C]^T$:

$$\text{softmax}(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^C e^{z_k}} \quad \text{for } j = 1, 2, \dots, C \quad (5)$$

Properties:

- All outputs in $(0, 1)$: $0 < \text{softmax}(\mathbf{z})_j < 1$
- Outputs sum to 1: $\sum_{j=1}^C \text{softmax}(\mathbf{z})_j = 1$
- Monotonic: larger $z_j \rightarrow$ larger output
- Translation invariant: $\text{softmax}(\mathbf{z} + c) = \text{softmax}(\mathbf{z})$

Use case: Multi-class classification output layer

SOFTMAX FUNCTION: NOTATION

Let's denote the softmax outputs as:

$$\hat{y}_j = \frac{e^{z_j}}{\sum_{k=1}^C e^{z_k}} \quad \text{where} \quad \hat{\mathbf{y}} = [\hat{y}_1, \hat{y}_2, \dots, \hat{y}_C]^T \quad (6)$$

Example: $\mathbf{z} = [1.0, 2.0, 3.0]^T$

$$\begin{aligned} \hat{y}_1 &= \frac{e^{1.0}}{e^{1.0} + e^{2.0} + e^{3.0}} = \frac{2.718}{2.718 + 7.389 + 20.086} \approx 0.090 \\ \hat{y}_2 &= \frac{e^{2.0}}{e^{1.0} + e^{2.0} + e^{3.0}} = \frac{7.389}{30.193} \approx 0.245 \\ \hat{y}_3 &= \frac{e^{3.0}}{e^{1.0} + e^{2.0} + e^{3.0}} = \frac{20.086}{30.193} \approx 0.665 \end{aligned}$$

Verify: $0.090 + 0.245 + 0.665 = 1.0 \checkmark$

SOFTMAX + CROSS-ENTROPY: BEAUTIFUL

Cross-entropy Loss: $\ell = - \sum_c y_c \log(\hat{y}_c)$ where true class is k (7)

Chain rule: $\frac{\partial \ell}{\partial z_j} = \sum_{i=1}^c \frac{\partial \ell}{\partial \hat{y}_i} \cdot \frac{\partial \hat{y}_i}{\partial z_j}$ (8)

$$\frac{\partial \ell}{\partial z_j} = -\frac{1}{\hat{y}_k} \cdot \frac{\partial \hat{y}_k}{\partial z_j} \quad \text{where} \quad \frac{\partial \ell}{\partial \hat{y}_i} = -\frac{y_i}{\hat{y}_i} \quad y_k = 1 \quad (9)$$

$$= -\frac{1}{\hat{y}_k} [\hat{y}_k (\delta_{kj} - \hat{y}_j)] \quad (10)$$

$$= -(\delta_{kj} - \hat{y}_j) \quad (11)$$

$$= \hat{y}_j - \delta_{kj} = \hat{y}_j - y_j \quad (12)$$

TANH FUNCTION: DEFINITION AND PROPERTIES

The hyperbolic tangent ($\tanh$) maps real numbers to the interval $(-1, 1)$.

Definition:

$$\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}} \quad (13)$$

Advantages over sigmoid:

- Zero-centered
- Stronger gradients (max = 1)
- Faster convergence

Properties:

- Output range: $(-1, 1)$
- Zero-centered: $\tanh(0) = 0$
- Odd function:
 $\tanh(-z) = -\tanh(z)$
- As $z \rightarrow \pm\infty$: $\tanh(z) \rightarrow \pm 1$

Key values:

- $\tanh(-1) \approx -0.76$
- $\tanh(0) = 0$
- $\tanh(1) \approx 0.76$

TANH FUNCTION: DERIVATIVE DERIVATION

Quotient Rule Method (where $\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$)

Let $u = e^z - e^{-z}$, $v = e^z + e^{-z}$, so $u' = e^z + e^{-z}$, $v' = e^z - e^{-z}$

$$\frac{d \tanh}{dz} = \frac{u'v - uv'}{v^2} = \frac{(e^z + e^{-z})^2 - (e^z - e^{-z})^2}{(e^z + e^{-z})^2} = \frac{4}{(e^z + e^{-z})^2} \quad (14)$$



Note: $(e^z + e^{-z})^2 - (e^z - e^{-z})^2 = (e^{2z} + 2 + e^{-2z}) - (e^{2z} - 2 + e^{-2z}) = 4$

TANH DERIVATIVE

$$\tanh'(z) = 1 - \tanh^2(z) = \frac{4}{(e^z + e^{-z})^2} \quad (15)$$

Beautiful property: Derivative computed using only the function value! Maximum at $z = 0$: $\tanh'(0) = 1$

TANH FUNCTION: VISUALIZATION



Key observations:

- Tanh is zero-centered (unlike sigmoid) - better for hidden layers
- Derivative peaks at $z = 0$ with value 1 (4× larger than sigmoid!)
- Derivative vanishes for large $|z|$ but slower than sigmoid
- Symmetric derivative (even function): $\tanh'(-z) = \tanh'(z)$

NUMERICAL PROBLEM: ACTIVATION FUNCTIONS

Problem: Given the input vector $z = \begin{bmatrix} 2.0 \\ -1.5 \\ 0.8 \\ -0.3 \end{bmatrix}$, calculate the output for each activation function. Also compute the derivatives for each function at the given input.

- 1 Sigmoid: $\sigma(z)$
- 2 Tanh: $\tanh(z)$
- 3 ReLU: $\text{ReLU}(z)$

SOLUTION: SIGMOID ACTIVATION

Sigmoid Function: $\sigma(z) = \frac{1}{1 + e^{-z}}$

$$\sigma(z) = \begin{bmatrix} \frac{1}{1+e^{-2.0}} \\ \frac{1}{1+e^{1.5}} \\ \frac{1}{1+e^{-0.8}} \\ \frac{1}{1+e^{0.3}} \end{bmatrix} = \begin{bmatrix} \frac{1}{1+0.135} = 0.881 \\ \frac{1}{1+4.482} = 0.182 \\ \frac{1}{1+0.449} = 0.690 \\ \frac{1}{1+1.350} = 0.426 \end{bmatrix} = \begin{bmatrix} 0.881 \\ 0.182 \\ 0.690 \\ 0.426 \end{bmatrix}$$

Sigmoid Derivative: $\sigma'(z) = \begin{bmatrix} 0.881(1 - 0.881) \\ 0.182(1 - 0.182) \\ 0.690(1 - 0.690) \\ 0.426(1 - 0.426) \end{bmatrix} = \begin{bmatrix} 0.105 \\ 0.149 \\ 0.214 \\ 0.245 \end{bmatrix}$

SOLUTION: TANH ACTIVATION

Tanh Function: $\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$

$$\tanh(z) = \begin{bmatrix} \tanh(2.0) \\ \tanh(-1.5) \\ \tanh(0.8) \\ \tanh(-0.3) \end{bmatrix} = \begin{bmatrix} \frac{7.389-0.135}{7.389+0.135} = 0.964 \\ \frac{0.223-4.482}{0.223+4.482} = -0.905 \\ \frac{2.226-0.449}{2.226+0.449} = 0.664 \\ \frac{0.741-1.350}{0.741+1.350} = -0.291 \end{bmatrix} = \begin{bmatrix} 0.964 \\ -0.905 \\ 0.664 \\ -0.291 \end{bmatrix}$$

Tanh Derivative: $\tanh'(z) = \begin{bmatrix} 1 - 0.964^2 \\ 1 - (-0.905)^2 \\ 1 - 0.664^2 \\ 1 - (-0.291)^2 \end{bmatrix} = \begin{bmatrix} 0.070 \\ 0.181 \\ 0.559 \\ 0.915 \end{bmatrix}$

SOLUTION: ReLU

ReLU Function: $\text{ReLU}(z) = \max(0, z)$

$$= \begin{bmatrix} \max(0, 2.0) \\ \max(0, -1.5) \\ \max(0, 0.8) \\ \max(0, -0.3) \end{bmatrix} = \begin{bmatrix} 2.0 \\ 0.0 \\ 0.8 \\ 0.0 \end{bmatrix}$$

ReLU Derivative: $\text{ReLU}'(z) = \begin{cases} 1 & \text{if } z > 0 \\ 0 & \text{if } z \leq 0 \end{cases}$

$$= \begin{bmatrix} 1 \\ 0 \\ 1 \\ 0 \end{bmatrix}$$

NUMERICAL PROBLEM: SOFTMAX FUNCTION

Problem: Given the logit vector z , compute:

- 1 The softmax output
- 2 The cross-entropy loss with true label y
- 3 The gradient of the loss with respect to the logits

$$z = \begin{bmatrix} 2.0 \\ 1.0 \\ 0.1 \\ -1.0 \end{bmatrix} \quad y = \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}$$

Cross-entropy Loss:

$$L = - \sum_{i=1}^n y_i \log(\text{softmax}(z_i))$$

SOLUTION: SOFTMAX COMPUTATION

Step 1: Compute softmax values First, for numerical stability, subtract the maximum:

$$z_{\text{stable}} = z - \max(z) = \begin{bmatrix} 2.0 \\ 1.0 \\ 0.1 \\ -1.0 \end{bmatrix} - 2.0 = \begin{bmatrix} 0.0 \\ -1.0 \\ -1.9 \\ -3.0 \end{bmatrix}$$

$$e^{z_{\text{stable}}} = \begin{bmatrix} e^{0.0} \\ e^{-1.0} \\ e^{-1.9} \\ e^{-3.0} \end{bmatrix} = \begin{bmatrix} 1.000 \\ 0.368 \\ 0.150 \\ 0.050 \end{bmatrix}$$

$$\sum_{j=1}^4 e^{z_j - \max(z)} = 1.000 + 0.368 + 0.150 + 0.050 = 1.568$$

SOLUTION: SOFTMAX COMPUTATION

$$\text{softmax}(z) = \frac{1}{1.568} \begin{bmatrix} 1.000 \\ 0.368 \\ 0.150 \\ 0.050 \end{bmatrix} = \begin{bmatrix} 0.638 \\ 0.235 \\ 0.096 \\ 0.032 \end{bmatrix}$$

$$\text{Verification: } = 0.638 + 0.235 + 0.096 + 0.032 = 1.001 \approx 1$$

SOLUTION: CROSS-ENTROPY LOSS AND GRADIENT

Step 2: Compute cross-entropy loss

$$L = - \sum_{i=1}^4 y_i \log(\text{softmax}(z_i))$$

$$y = \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix} \quad \text{only the second term contributes}$$

$$L = -y_2 \log(\text{softmax}(z_2)) = -1 \times \log(0.235) = -\log(0.235) = 1.447$$

SOLUTION: CROSS-ENTROPY LOSS AND GRADIENT

Step 3: Compute gradient For cross-entropy loss with softmax, the gradient is:

$$\frac{\partial L}{\partial z_i} = \text{softmax}(z_i) - y_i$$
$$\frac{\partial L}{\partial \mathbf{z}} = \begin{bmatrix} 0.638 \\ 0.235 \\ 0.096 \\ 0.032 \end{bmatrix} - \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix} = \begin{bmatrix} 0.638 \\ -0.765 \\ 0.096 \\ 0.032 \end{bmatrix}$$

This result shows that the gradient is simply the difference between predicted and true probabilities!

TABLE OF CONTENTS

- 1 COMPLETE NEURAL NETWORK TRAINING
- 2 GRADIENT FLOW
- 3 DESIGN A NEURAL NETWORK
- 4 ACTIVATION FUNCTIONS
- 5 LOSS FUNCTIONS AND THEIR DERIVATIVES
- 6 COMPARISON AND SUMMARY



OVERVIEW: LOSS FUNCTIONS

Loss functions measure the discrepancy between predicted and actual values.

We will include:

- 1 Mean Squared Error (MSE) Loss
- 2 Mean Absolute Error (MAE) Loss
- 3 Binary Cross-Entropy (BCE) Loss
- 4 Cross-Entropy Loss

Key Goal: Understand how to compute gradients for backpropagation

MSE LOSS: DEFINITION

Mean Squared Error (MSE) is commonly used for regression tasks.

$$\ell(\hat{y}, y) = \frac{1}{2}(\hat{y} - y)^2 \quad (16)$$

Properties:

- Always non-negative: $\ell \geq 0$
- Convex function (for linear models)
- Differentiable everywhere
- Penalizes large errors quadratically
- Factor of $\frac{1}{2}$ simplifies derivative

MSE LOSS: DERIVATIVE DERIVATION

Goal: Find $\frac{\partial \ell}{\partial \hat{y}}$ where $\ell = \frac{1}{2}(\hat{y} - y)^2$

Let $u = \hat{y} - y$, so $\ell = \frac{1}{2}u^2$

$$\frac{\partial \ell}{\partial \hat{y}} = \frac{\partial \ell}{\partial u} \cdot \frac{\partial u}{\partial \hat{y}} \quad (\text{chain rule}) \quad (17)$$

$$= \left(\frac{1}{2} \cdot 2u \right) \cdot (1) \quad (\text{compute derivatives}) \quad (18)$$

$$= u = \hat{y} - y \quad (\text{substitute back}) \quad (19)$$

MSE Loss Derivative:

$$\boxed{\frac{\partial \ell}{\partial \hat{y}} = \hat{y} - y} \quad (20)$$

MSE LOSS: NUMERICAL EXAMPLE

Example 1: $y = 3, \hat{y} = 5$

$$\ell = \frac{1}{2}(5 - 3)^2 = \frac{1}{2}(2)^2 = 2$$
$$\frac{\partial \ell}{\partial \hat{y}} = 5 - 3 = 2$$

Example 2: $y = 8, \hat{y} = 5$

$$\ell = \frac{1}{2}(5 - 8)^2 = \frac{1}{2}(-3)^2 = 4.5$$
$$\frac{\partial \ell}{\partial \hat{y}} = 5 - 8 = -3$$

Note: Gradient is negative, telling us to *increase* $\hat{y}$

MAE Loss: DEFINITION

Mean Absolute Error (MAE) is more robust to outliers than MSE.

$$\ell(\hat{y}, y) = |\hat{y} - y| \quad (21)$$

Properties:

- Always non-negative: $\ell \geq 0$
- Less sensitive to outliers than MSE
- Not differentiable at $\hat{y} = y$
- Linear penalty for errors

MAE LOSS: DERIVATIVE DERIVATION

Goal: Find $\frac{\partial \ell}{\partial \hat{y}}$ where $\ell = |\hat{y} - y|$

Recall: Derivative of absolute value function

$$\frac{d|x|}{dx} = \begin{cases} 1 & \text{if } x > 0 \\ -1 & \text{if } x < 0 \\ \text{undefined} & \text{if } x = 0 \end{cases} \quad (22)$$

Apply to $\ell = |\hat{y} - y|$ using chain rule:

$$\frac{\partial \ell}{\partial \hat{y}} = \frac{\partial}{\partial \hat{y}} |\hat{y} - y| = \begin{cases} 1 & \text{if } \hat{y} > y \\ -1 & \text{if } \hat{y} < y \\ 0 & \text{if } \hat{y} = y \text{ (convention)} \end{cases} = \text{sign}(\hat{y} - y) \quad (23)$$

MAE Loss Derivative:

$$\boxed{\frac{\partial \ell}{\partial \hat{y}} = \text{sign}(\hat{y} - y)} \quad (24)$$

MAE LOSS: NUMERICAL EXAMPLES

Example 1: $y = 3, \hat{y} = 5$

$$\ell = |5 - 3| = 2$$

$$\frac{\partial \ell}{\partial \hat{y}} = \text{sign}(5 - 3) = \text{sign}(2) = +1$$

Example 2: $y = 8, \hat{y} = 5$

$$\ell = |5 - 8| = 3$$

$$\frac{\partial \ell}{\partial \hat{y}} = \text{sign}(5 - 8) = \text{sign}(-3) = -1$$

Comparison with MSE:

- MAE gradient is constant: ± 1
- MSE gradient varies with error magnitude: $\hat{y} - y$

BCE LOSS: DEFINITION

Binary Cross-Entropy (BCE) is used for binary classification where $y \in \{0, 1\}$ and $\hat{y} \in (0, 1)$.

For a single example:

$$\ell(\hat{y}, y) = -[y \log(\hat{y}) + (1 - y) \log(1 - \hat{y})] \quad (25)$$

Properties:

- Measures divergence between predicted and true probability distributions
- Always non-negative: $\ell \geq 0$
- Heavily penalizes confident wrong predictions

BCE LOSS: UNDERSTANDING THE FORMULA

Why this form?

Since $y \in \{0, 1\}$, the loss simplifies based on the true label:

Case 1: When $y = 1$

$$\ell = -[1 \cdot \log(\hat{y}) + (1 - 1) \log(1 - \hat{y})] \quad (26)$$

$$= -\log(\hat{y}) \quad (27)$$

Case 2: When $y = 0$

$$\ell = -[0 \cdot \log(\hat{y}) + (1 - 0) \log(1 - \hat{y})] \quad (28)$$

$$= -\log(1 - \hat{y}) \quad (29)$$

Intuition:

- If $y = 1$, want $\hat{y} \rightarrow 1$ to minimize $-\log(\hat{y})$
- If $y = 0$, want $\hat{y} \rightarrow 0$ to minimize $-\log(1 - \hat{y})$

BCE LOSS: DERIVATIVE DERIVATION

Goal: Find $\frac{\partial \ell}{\partial \hat{y}}$ where $\ell = -[y \log(\hat{y}) + (1 - y) \log(1 - \hat{y})]$

Differentiate both terms (using $\frac{d}{dx} \log(x) = \frac{1}{x}$ and chain rule):

$$\frac{\partial \ell}{\partial \hat{y}} = -\frac{\partial}{\partial \hat{y}} [y \log(\hat{y})] - \frac{\partial}{\partial \hat{y}} [(1 - y) \log(1 - \hat{y})] \quad (30)$$

$$= -\frac{y}{\hat{y}} - (1 - y) \cdot \frac{1}{1 - \hat{y}} \cdot (-1) = -\frac{y}{\hat{y}} + \frac{1 - y}{1 - \hat{y}} \quad (31)$$

Simplify with common denominator:

$$\frac{\partial \ell}{\partial \hat{y}} = \frac{-y(1 - \hat{y}) + (1 - y)\hat{y}}{\hat{y}(1 - \hat{y})} = \frac{-y + y\hat{y} + \hat{y} - y\hat{y}}{\hat{y}(1 - \hat{y})} = \frac{\hat{y} - y}{\hat{y}(1 - \hat{y})} \quad (32)$$

BCE Loss Derivative:

$$\boxed{\frac{\partial \ell}{\partial \hat{y}} = \frac{\hat{y} - y}{\hat{y}(1 - \hat{y})}} \quad (33)$$

BCE LOSS: DERIVATIVE EXAMPLES

Example 1: $y = 1, \hat{y} = 0.9$

$$\ell = -\log(0.9) \approx 0.105$$

$$\frac{\partial \ell}{\partial \hat{y}} = \frac{0.9 - 1}{0.9(1 - 0.9)} = \frac{-0.1}{0.09} \approx -1.11$$

Example 2: $y = 1, \hat{y} = 0.1$

$$\ell = -\log(0.1) \approx 2.303$$

$$\frac{\partial \ell}{\partial \hat{y}} = \frac{0.1 - 1}{0.1(1 - 0.1)} = \frac{-0.9}{0.09} = -10$$

Example 3: $y = 0, \hat{y} = 0.9$

$$\ell = -\log(0.1) \approx 2.303$$

$$\frac{\partial \ell}{\partial \hat{y}} = \frac{0.9 - 0}{0.9(1 - 0.9)} = \frac{0.9}{0.09} = 10$$

CROSS-ENTROPY LOSS: DEFINITION

Cross-Entropy Loss generalizes BCE to multi-class classification where y is one-hot encoded and $\hat{\mathbf{y}}$ is a probability distribution.

For C classes and single example:

$$\ell(\hat{\mathbf{y}}, \mathbf{y}) = - \sum_{c=1}^C y_c \log(\hat{y}_c) \quad (34)$$

where:

- $\mathbf{y} = [y_1, y_2, \dots, y_C]^T$ is one-hot encoded (one $y_c = 1$, rest are 0)
- $\hat{\mathbf{y}} = [\hat{y}_1, \hat{y}_2, \dots, \hat{y}_C]^T$ are predicted probabilities
- $\sum_{c=1}^C \hat{y}_c = 1$ (output of softmax)

CROSS-ENTROPY LOSS: DERIVATIVE DERIVATION

Goal: Find $\frac{\partial \ell}{\partial \hat{y}_j}$ for each class j , where $\ell = -\sum_{c=1}^C y_c \log(\hat{y}_c)$

Step 1: Take derivative with respect to $\hat{y}_j$

$$\frac{\partial \ell}{\partial \hat{y}_j} = -\frac{\partial}{\partial \hat{y}_j} \left[\sum_{c=1}^C y_c \log(\hat{y}_c) \right] \quad (35)$$

Step 2: Only the $c = j$ term depends on $\hat{y}_j$

$$\frac{\partial \ell}{\partial \hat{y}_j} = -\frac{\partial}{\partial \hat{y}_j} [y_j \log(\hat{y}_j)] = -y_j \cdot \frac{1}{\hat{y}_j} \quad (36)$$

Cross-Entropy Loss

$$\boxed{\frac{\partial \ell}{\partial \hat{y}_j} = -\frac{y_j}{\hat{y}_j}} \quad (37)$$

CROSS-ENTROPY LOSS: GRADIENT VECTOR

The gradient with respect to all outputs forms a vector:

$$\nabla_{\hat{\mathbf{y}}} \ell = \begin{bmatrix} \frac{\partial \ell}{\partial \hat{y}_1} \\ \frac{\partial \ell}{\partial \hat{y}_2} \\ \vdots \\ \frac{\partial \ell}{\partial \hat{y}_C} \end{bmatrix} = \begin{bmatrix} -\frac{y_1}{\hat{y}_1} \\ -\frac{y_2}{\hat{y}_2} \\ \vdots \\ -\frac{y_C}{\hat{y}_C} \end{bmatrix} \quad (38)$$

Since $\mathbf{y}$ is one-hot (only $y_k = 1$):

- For the true class k : $\frac{\partial \ell}{\partial \hat{y}_k} = -\frac{1}{\hat{y}_k}$
- For all other classes $j \neq k$: $\frac{\partial \ell}{\partial \hat{y}_j} = 0$

Important: This is the gradient w.r.t. softmax outputs. When combined with softmax derivative, we get a simpler form!

TABLE OF CONTENTS

- 1 COMPLETE NEURAL NETWORK TRAINING
- 2 GRADIENT FLOW
- 3 DESIGN A NEURAL NETWORK
- 4 ACTIVATION FUNCTIONS
- 5 LOSS FUNCTIONS AND THEIR DERIVATIVES
- 6 COMPARISON AND SUMMARY



ACTIVATION FUNCTIONS: COMPARISON

Property	Sigmoid	Tanh	Softmax
Formula	$\frac{1}{1+e^{-z}}$	$\frac{e^z - e^{-z}}{e^z + e^{-z}}$	$\frac{e^{z_j}}{\sum_k e^{z_k}}$
Range	$(0, 1)$	$(-1, 1)$	$(0, 1)$
Zero-centered?	No	Yes	No
Derivative	$\sigma(1 - \sigma)$	$1 - \tanh^2$	Complex
Max Derivative	0.25	1.0	Varies
Output Sum	N/A	N/A	1.0
Use Case	Binary class	Hidden layers	Multi-class

Key Advantages:

- **Sigmoid:** Output is probability for binary classification
- **Tanh:** Zero-centered, stronger gradients
- **Softmax:** Outputs form probability distribution

LOSS FUNCTIONS: COMPARISON

Loss	Formula	Derivative
MSE	$\frac{1}{2}(\hat{y} - y)^2$	$\hat{y} - y$
MAE	$ \hat{y} - y $	$\text{sign}(\hat{y} - y)$
BCE	$-[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$	$\frac{\hat{y} - y}{\hat{y}(1 - \hat{y})}$
Cross-Entropy	$-\sum_c y_c \log \hat{y}_c$	$-\frac{y_j}{\hat{y}_j}$

Use Cases:

- **MSE:** Regression tasks
- **MAE:** Regression with outliers
- **BCE:** Binary classification ($y \in \{0, 1\}$)
- **Cross-Entropy:** Multi-class classification

Special case: Softmax + Cross-Entropy gradient: $\hat{y}_j - y_j$ (very clean!)

SUMMARY: KEY TAKEAWAYS

Loss Functions:

- MSE penalizes errors quadratically, derivative is linear
- MAE is robust to outliers, derivative is constant ± 1
- BCE for binary classification, heavily penalizes confident errors
- Cross-Entropy for multi-class, only cares about correct class probability

Activation Functions:


- Sigmoid squashes to $(0, 1)$, useful for probabilities
- Tanh squashes to $(-1, 1)$, zero-centered is better for hidden layers
- Softmax converts logits to probability distribution

Practical Tips:

- Regression: Use MSE or MAE loss
- Binary classification: Use sigmoid + BCE
- Multi-class: Use softmax + cross-entropy

TRY

- Neural Network Playground: playground.tensorflow.org



Thank You!